

# Project organization

David Benkeser, PhD MPH

Emory University

Department of Biostatistics and Bioinformatics

INFO550

 [bit.ly/info550](https://bit.ly/info550)

 Additional reading

*File organization and naming are powerful weapons against chaos.*

Jenny Bryan

# Basic principles

- **Put everything in one version-controlled directory.**
- Develop your own system.
- Be consistent, but look for ways to improve.
  - naming conventions, file structure, `make` structure
- Raw data are sacred. Keep them separate from everything else.
- Separate code and data.
- Use `make` files and/or READMEs to document dependencies.
- No spaces in file names.
- Use meaningful file names.
- Use YYYY-MM-DD date formatting.
- **No absolute paths.**
- **Use a package management system.**

*You mostly collaborate with yourself, and me-from-two-months-ago never responds to email.*

Karen Cranston

# What to organize?

It is probably useful to have a system for organizing:

- data analysis projects;
- first-author papers;
- talks.

The systems should adhere to the same general principles, but different requirements may necessitate different structures.

**Think about organization of a project from the outset!**

# Collaborative projects

Collaborative projects present a greater challenge.

- Not everyone is comfortable with LaTeX or git or ...

I don't have a great solution for this.

- Google drive/Word online helps to a certain extent, but you lose in other areas (reference management, math typesetting)
- [Overleaf](#) has gotten much better for LaTeX

Some advice:

- Address organization from the outset.
- Ideally, bring people on board to your (version controlled, reproducible) system.
- Keep open lines of communication (especially if using GitHub)

# Example data analysis project

```
analysis/  
  raw_data/  
  data/  
  R/  
    R/00_clean_data.R  
    R/01_fit_models.R  
    R/02_make_figures.R  
    R/03_summarize_results.R  
    R/04_report.Rmd  
  figs/  
  sandbox/  
    sandbox/exploratory.R  
  ref_papers/  
  Makefile  
  README.md  
  renv
```

# Example paper organization

```
paper/  
  analysis/  
    analysis/README.md  
    analysis/00_clean_data.R  
    analysis/01_fit_models.R  
    analysis/02_make_figures.R  
    analysis/sandbox  
  sim/  
    sim/README.md  
    sim/helper_functions.R  
    sim/sim_script.R  
    sim/run_sim_script.sh  
    sim/sandbox  
  figs/  
  notes/  
  ref_papers/  
  submitted/  
  revision/  
  final/  
  README.md  
  Makefile  
  my_paper.tex  
  my_refs.bib
```



# Example class organization

See the `class` GitHub repository!

# Organizing data

Raw data are sacred... but may be a mess.

- You'll be surprised (and disheartened) by how many color-coded excel sheets you'll get in your life.

Tempting to edit raw data by hand. **Don't!**

- Everything scripted!

Use meta-data files to describe raw and cleaned data.

- structure as data (e.g., `.csv` so easy to read)

# Organizing data

Hadley Wickham defined the notion of tidy data.

- Each variable forms a column.
- Each observation forms a row.
- Each observational unit forms a table.

ptid	day	age	drug	out
1	1	28	0	0
1	2	28	0	1
2	1	65	0	0
2	2	65	1	1
3	1	34	0	0
3	2	34	–	1

# Exploring data

One of the first things we'll often do is open the data and start poking around.

- Could be informal, "getting to know you."
- Could be more formal, "see if anything looks interesting."

This is often done in an ad-hoc way:

- entering commands directly into R;
- making and saving plots "by hand";
- etc...

**Slow down and document.**

- Your future self will thank you!

# Exploring data

Other helpful ideas for formalizing exploratory data analysis:

- `.Rhistory` files
  - all the commands used in an R session
- Informal `.Rmd` documents.
  - easy way to organize code/comments into readable format
- `save` intermediate objects and workspaces
  - and document what they contain!
- `knitr::spin`
  - writing `.R` scripts with rendered-able comments

# The here package

## No absolute paths.

- Absolute paths are the enemy of project reproducibility.

For [R](#) projects, the [here](#) package provides a simple way to use relative file paths.

- Read [Jenny Bryan and James Hester's chapter](#) on project-oriented work-flows.

The use of [here](#) is dead-simple and best illustrated by example.

# The here package

Consider this simple project structure.

```
my_project/  
  data/  
    my_data.csv  
  output/  
  R/  
    R/my_analysis.R  
  Rmd/  
    Rmd/my_report.Rmd
```

Here, the folder `my_project` is the **root directory**.

- Where `.git` lives
- All file paths should be **relative** to `my_project`!

# The here package

Each R script or Rmd report, should contain a call to `here::i_am('path/to/this/file')` at the top.

- `path/to/this/file` should be replaced with the path **relative** to the project's **root directory**.
- `here::i_am` means use function `i_am` from `here` package.

For example, the file `R/my_analysis.R` might look like this.

```
# include at top of script
here::i_am('R/my_analysis.R')

# now add all your great R code...
```



# The here package

Rmd/my\_report.Rmd should include an R chunk that calls `i_am`.

```
---  
output: html_document  
---  
  
{r}  
here::i_am('Rmd/my_report.Rmd')  
  
<!-- Now the rest of your Rmd code -->
```

# The here package

The call to `i_am` establishes the **root directory**.

- Subsequent file paths can be made using the `here` function.

For example, `my_analysis.R` might look like this.

```
# include at top of script
here::i_am('R/my_analysis.R')

# load data
my_data <- read.csv(here::here('data', 'my_data.csv'))

# do some analysis to get results

# save results
save(my_results, file = here::here('output', 'my_results.RData'))
```

# The `renv` package

## Use a package management system.

To increase reproducibility of a project, we must keep track of what packages are used.

- Want to avoid chasing down 100 errors like this:

```
Error in library(ggplot2) : there is no package called  
'ggplot2'
```

The `renv` package is useful to this end.

# The renv package

Download the `example_project` folder.

```
example_project/  
  Makefile  
  figs/  
  R/  
    R/barchart.R  
  Rmd/  
    Rmd/report.Rmd
```

If you scan the code, you will see we need **three R packages**.

- `here`, `wesanderson`, `knitr`, and `rmarkdown`

# The `renv` package

Open an R session from the `example_project` folder.

- Install the `renv` package if needed.

Initialize the project by running the following command.

```
renv::init()
```

You will see a lot of output. What just happened??

# renv.lock file

You should now see **lockfile**, `example_project/renv.lock`.

```
{
  "R": {
    "Version": "4.0.3",
    "Repositories": [
      {
        "Name": "CRAN",
        "URL": "https://cloud.r-project.org"
      }
    ]
  },
  "Packages": {
    "base64enc": {
      "Package": "base64enc",
      "Version": "0.1-3",
      "Source": "Repository",
      "Repository": "CRAN",
      "Hash": "543776ae6848fde2f48ff3816d0628bc"
    },
    ...
  }
}
```

# renv.lock file

This **lockfile** records all of the information about packages needed by your project.

- Version of package, where was it installed from, etc...

How does it know?

- **renv** scans all files in your project directory.
- Looks for **library**, **require**, or **package::function**.

# .Rprofile

We also see `example_project/.Rprofile`, containing one line.

```
source("renv/activate.R")
```

We have not talked about `.Rprofile`'s because they are generally antithetical to reproducible research.

When `R` starts, it searches for `.Rprofile` and runs what it finds.

- Use this to change various options, always load packages, etc...

In this case, whenever we start `R` in `example_project`, we **activate** our `R` environment.

- Telling `R` where packages for this project are saved, etc...
- Details are not too important.



# renv folder

You will also see a folder `example_project/renv/`.

- This folder contains your **project library**.

`renv` tries to be clever about installing packages.

- Already have a package installed elsewhere? `renv` will link to it.
- Otherwise, package is installed in `renv/library`.

Note that `renv/.gitignore` ensures packages not put under version control.

# Activating renv

`renv` will automatically be active in any R session that is run from the `example_project` directory.

- Recall the presence and function of `.Rprofile`.

To activate in an R session run from elsewhere:

```
renv::activate('path/to/renv').
```

For using `renv` with R Markdown projects, this is important!

- Either need to `activate` in a code chunk.
- Or use `knitr::opts_knit$set(root.dir = here::here())` (assuming the project root contains your `renv` project library).

# Collaborating with `renv`

A typical `renv` collaborative workflow on GitHub:

- **User A** initializes the lockfile using `renv::init()`.
- **User A** commits `renv.lock`, `.Rprofile`, and `renv/activate.R` and pushes to GitHub.
- **User B** pulls from GitHub, opens `R`, and uses `renv::restore()` to synchronize their local project directory.
- **User B** adds new packages to code, uses `renv::snapshot()` to record changes to `renv.lock`
- **User B** commits `renv.lock` and pushes to GitHub.
- **User A** pulls from GitHub, opens `R`, and uses `renv::restore()` to synchronize their local project directory.
- ...

# Breakout exercise

With a partner, choose **User A** and **User B**.

## User A

- initialize the lockfile for `example_project`
  - open R in `example_project` and run `renv::init()`
- use `git init` to initialize version control of `example_project`
- create a GitHub repository for `example_project`
- add the GitHub repository as a remote to your local repository
  - `git remote add origin`  
`https://github.com/usera/example_project.git`
  - `git remote add origin`  
`git@github.com:usera/example_project.git`
- commit all files locally and push to repository

# Breakout exercise

## User B

- create a fork of **User A**'s repository
- `git clone` the fork to your local machine
- `cd` into `example_project` and open R
- run `renv::restore()` to synchronize package library
- confirm that you can build the report
  - run `make report` from `example_project` directory
  - open `Rmd/report.html` to confirm correct build

# Breakout exercise

**User B** now wants to change the colors of the graph

- open R in `example_project` directory
- run `renv::remove('wesanderson')` to remove the `wesanderson` package from the lockfile
- replace lines 4–5 of `barchart.R` with the following and save

```
library(RColorBrewer)
colors <- brewer.pal(3, "Dark2")
```

- open R in `example_project` and run `renv::status()`
- if prompted, run `install.packages("RColorBrewer")`
- run `renv::snapshot()` to add `RColorBrewer` to the lockfile
- commit changes and push to your fork
- submit a PR to **User A's** repository

# Breakout exercise

## User A

- add a remote linking to user B's repository
  - `git remote add userb`  
`https://github.com/userb/example_project.git`
- fetch **User B's** master branch
  - `git fetch userb master`
- checkout **User B's** master branch
  - `git checkout remotes/userb/master`

# Breakout exercise

## User A

- `cd` into `example_project` and open R
- run `renv::restore()` to synchronize your package library
- confirm that you can build the report
  - run `make report` from `example_project` directory
  - open `Rmd/report.html` to confirm correct build



# Breakout exercise

## User A

- if the report builds correctly, create a local branch named `userb` from **User B's** `master` branch
  - `git checkout -b userb`
- merge **User B's** `master` into your `master`
  - `git checkout master`
  - `git merge userb`
- close the PR by pushing to GitHub
  - `git push origin master`

# Breakout exercise

User A has changed their mind about the colors!

- open R in `example_project` directory
- run `renv::remove('RColorBrewer')` to remove `RColorBrewer` from the lockfile
- replace lines 4–5 of `barchart.R` with the following and save

```
colors <- c("red", "blue", "green")
```

- open R and run `renv::status()` and `renv::snapshot()`
- confirm that you can build the report
  - run `make report` from `example_project` directory
  - open `Rmd/report.html` to confirm correct build
- if the report builds correctly, `commit` and `push`

# Breakout exercise

## User B

- add a remote linking to user A's repository
  - `git remote add usera https://github.com/usera/example_project.git`
- fetch **User A's** master branch
  - `git fetch usera master`
- checkout **User A's** master branch
  - `git checkout remotes/usera/master`

# Breakout exercise

## User B

- `cd` into `example_project` and open R
- run `renv::status()` and, if needed, `renv::restore()` to synchronize your package library
- confirm that you can build the report
  - run `make report` from `example_project` directory
  - open `Rmd/report.html` to confirm correct build

# Breakout exercise

## User B

- if the report builds correctly, create a local branch named `usera` from **User A's** `master` branch
  - `git checkout -b usera`
- merge **User A's** `master` into your `master`
  - `git checkout master`
  - `git merge usera`
- update your fork of the repository by pushing to GitHub
  - `git push origin master`